

draw2pix: Generative Networks for Bad Artists

Cedrick Argueta
Stanford University
450 Serra Mall, Stanford, CA 94305
cedrick@cs.stanford.edu

Kevin Wang
Stanford University
450 Serra Mall, Stanford, CA 94305
kwang98@stanford.edu

1. Introduction

We are investigating unpaired image-to-image translation with generative adversarial networks. This work was demonstrated in [5]. Work done in the same lab [2] demonstrates how we can use paired images in training to create an image-to-image mapping from one domain to another. The specific image-to-image task that we’re aiming to implement is this, which demonstrates a couple examples of sketch-to-photo translations trained using the pix2pix model. We aim to do something similar, using CycleGAN to instead train on unpaired images. The relaxation that unpaired training gives us allows for easier creation of datasets and combinations of domains – we can simply swap a whole domain and retrain rather than find image-to-image pairs for that specific translation task. Our project differs from [6] in that we aren’t doing multi-modal image translation, just from one domain to another. We are essentially implementing the demo from [2] with an unpaired dataset and architecture matching that of [5]

2. Related Work

describe some of these, how they relate
The original Generative Adversarial Network (GAN) proposed by Goodfellow et al. [1] proposes a network that learns an approximation of a distribution of data p_X . It does so by training two networks - a generator network and a discriminator network. These networks are trained adversarially, i.e. the discriminator’s objective is to discern between real images and the generator’s images, and the generator’s objective is to generate images that fool the discriminator.

Generative networks are not limited to simple distributions. Work done by Liu et al. [3] demonstrated the ability of generative networks to learn a joint distribution between variables, through the use of multiple generator and discriminator networks. This can be used to generate pairs $x_1 \sim p_{X_1}, x_2 \sim p_{X_2}$ of images that belong to the joint distribution just by using the marginal distributions.

Zhu et al. [5] goes a step further by enforcing cycle consistency between pairs created, i.e. by ensuring that gener-

ated images can be run through another generator to create the paired image. CycleGAN effectively learns a mapping function $G : X \rightarrow Y$ for images in domains X and Y .

The networks discussed earlier are performing unsupervised image-to-image translation in that no labels are given to pair images from domain X and domain Y together. When these labels are given, supervised methods like those described by Isola et al. [2] are possible through the use of conditional GANs.

3. Problem Statement

The problem we are addressing is image-to-image translation from an input domain X of hand-drawn tree sketches to the output domain Y of realistic photographs of trees. Specifically, we are working with unpaired training data, meaning there is no mapping between sketches $x_i \in X$ and target photos $y_j \in Y$. Often, finding paired data for the image-to-image translation task is difficult. Moreover, using unpaired image data allows for independent collection of input and output training sets. This motivates the problem that we are attempting to solve.

4. Technical Approach

We aim to use CycleGAN [5] for unpaired image-to-image translation. More precisely, we will train networks to learn functions $G : X \rightarrow Y$ and $F : Y \rightarrow X$. This will allow us to take an image $x_1 \in X$ and generate its pair $y_1 \in Y$, and vice versa. We let the true distributions of the domains be p_X and p_Y . Just as in [1], we will use discriminators D_X and D_Y as adversarial networks to the generators, which we will train to distinguish between elements in X and p_X and Y and p_Y , respectively. CycleGAN uses two losses per generator function. The first loss is adversarial, where the loss function for G and its discriminator D_Y is:

$$\mathcal{L}_{GAN}(G, D_Y, X, Y) = \mathbb{E}_{y \sim p_Y} [\log D_Y(y)] + \mathbb{E}_{x \sim p_X} [\log (1 - D_Y(G(x)))] \quad (1)$$

and the objective is $\min_G \max_{D_Y} \mathcal{L}_{GAN}$. We can apply this without loss of generality to F and its discriminator D_X .

Because it is possible for a network to map an input image x to multiple images in Y , we use a cycle consistency loss to guarantee that there is a unique mapping. Specifically, we wish to enforce that for images x and y $F(G(x)) \approx x$ (forward cycle consistency) and $G(F(y)) \approx y$ (backward cycle consistency). The loss used is:

$$\mathcal{L}_{cyc} = \mathbb{E}_{x \sim p_X} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p_Y} [\|G(F(y)) - y\|_1] \quad (2)$$

Then the full loss function is:

$$\begin{aligned} \mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{GAN}(G, D_Y, X, Y) \\ & + \mathcal{L}_{GAN}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{cyc}(G, F) \end{aligned} \quad (3)$$

where λ is a parameter that determines the importance of cyclic consistency loss vs. adversarial loss. Minimizing this loss through

$$G^*, F^* = \arg \min_{G, F} \max_{D_X, D_Y} \mathcal{L}(G, F, D_X, D_Y) \quad (4)$$

yields G^* and F^* , our optimal mapping functions.

5. Dataset

We will be using a subset of the Sketchy dataset, as described in [4]. The dataset consists of a crowd-sourced collection of sketch-photo pairs. The total collection consists of 75,471 sketches of 12,500 photos (125 object categories). We will be working specifically with the 'tree' subset, which consists of 100 photos and 534 sketches. Although there is a naturally defined pairing between photos and sketches, we will be ignoring this mapping in our implementation of unpaired translation. Currently, we are preprocessing both domains in our baseline by zero-centering and scaling pixel values to the $[-1, 1]$ range. We are planning on normalizing data for both domains in our implementation of CycleGAN as well.

6. Baseline Results

describe cogan method, results

References

- [1] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial nets. In *Advances in neural information processing systems*, pages 2672–2680, 2014.
- [2] P. Isola, J. Zhu, T. Zhou, and A. A. Efros. Image-to-image translation with conditional adversarial networks. *CoRR*, abs/1611.07004, 2016.
- [3] M. Liu and O. Tuzel. Coupled generative adversarial networks. *CoRR*, abs/1606.07536, 2016.
- [4] P. Sangkloy, N. Burnell, C. Ham, and J. Hays. The sketchy database: learning to retrieve badly drawn bunnies. *ACM Transactions on Graphics (TOG)*, 35(4):119, 2016.
- [5] J. Zhu, T. Park, P. Isola, and A. A. Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. *CoRR*, abs/1703.10593, 2017.
- [6] J. Zhu, R. Zhang, D. Pathak, T. Darrell, A. A. Efros, O. Wang, and E. Shechtman. Toward multimodal image-to-image translation. *CoRR*, abs/1711.11586, 2017.